

# Exploratory data analysis with pandas in python, part 1

**Andras Zsom, Lead Data Scientist at CCV**

`github.com/brown-ccv/dscov_data_science`

## Series of talks by yours truly

- DATA1030 - Hands-on Data Science
  - how to develop ML pipelines for tabular data from scratch
- walk through a modified version of the course material
- today:
  - how to manipulate tabular data with pandas
- down the road:
  - how to visualize your data
  - ML overview
  - discuss each step of an ML pipeline

## Data transformations: pandas data frames

### By the end of this presentation, you will be able to

- read in csv, excel, and sql data into a pandas data frame
- filter rows in various ways
- select columns
- merge and append data frames

### Some notes and advice

- **ALWAYS READ THE HELP OF THE METHODS/FUNCTIONS YOU USE!**
- stackoverflow is your friend, use it! <https://stackoverflow.com/> (<https://stackoverflow.com/>)

# Pandas

- data are often distributed over multiple files/databases (e.g., csv and excel files, sql databases)
- each file/database is read into a pandas dataframe
- you often need to filter dataframes (select specific rows/columns based on index or condition)
- pandas dataframes can be merged and appended

## Data transformations: pandas data frames

By the end of this talk, you will be able to

- **read in csv, excel, and sql data into a pandas data frame**
- filter rows in various ways
- select columns
- merge and append data frames

```
In [1]: # how to read in a database into a dataframe and basic dataframe structure
import pandas as pd

# load data from a csv file
df = pd.read_csv('data/adult_data.csv') # there are also pd.read_excel(), and pd.read_sql()

#print(df)
#print(df.head()) # by default, shows the first five rows but check help (df.head) to specify the number of rows to show
#print(df.shape) # the shape of your dataframe (number of rows, number of columns)
#print(df.shape[0]) # number of rows
print(df.shape[1]) # number of columns
```

## Packages

A package is a collection of classes and functions.

- a dataframe (`pd.DataFrame()`) is a pandas class
  - a class is the blueprint of how the data should be organized
  - classes have methods which can perform operations on the data (e.g., `.head()`, `.shape`)
- `df` is an object, an instance of the class.
  - we put data into the class
  - methods are attached to objects
    - you cannot call `pd.head()`, you can only call `df.head()`
- `read_csv` is a function
  - functions are called from the package
  - you cannot call `df.read_csv`, you can only call `pd.read_csv()`

## DataFrame structure: both rows and columns are indexed!

- index column, no name
  - contains the row names
  - by default, index is a range object from 0 to number of rows - 1
  - any column can be turned into an index, so indices can be non-number, and also non-unique. more on this later.
- columns with column names on top

## Always print your dataframe to check if it looks ok!

### Most common reasons it might not look ok:

- the first row is not the column name
  - there are rows above the column names that need to be skipped
  - there is no column name but by default, pandas assumes the first row is the column name. as a result, the values of the first row end up as column names.
- character encoding is off
- separator is not comma but some other character

```
In [2]: # check the help to find the solution  
help(pd.read_csv)
```

Help on function read\_csv in module pandas.io.parsers:

```
read_csv(filepath_or_buffer:Union[str, pathlib.Path, IO[~AnyStr]], sep
=' ', delimiter=None, header='infer', names=None, index_col=None, useco
ls=None, squeeze=False, prefix=None, mangle_dupe_cols=True, dtype=None,
engine=None, converters=None, true_values=None, false_values=None, skip
initialspace=False, skiprows=None, skipfooter=0, nrows=None, na_values=
None, keep_default_na=True, na_filter=True, verbose=False, skip_blank_l
ines=True, parse_dates=False, infer_datetime_format=False, keep_date_co
l=False, date_parser=None, dayfirst=False, cache_dates=True, iterator=F
alse, chunksize=None, compression='infer', thousands=None, decimal=
b'.' , lineterminator=None, quotechar='"', quoting=0, doublequote=True,
escapechar=None, comment=None, encoding=None, dialect=None, error_bad_l
ines=True, warn_bad_lines=True, delim_whitespace=False, low_memory=Tru
e, memory_map=False, float_precision=None)
```

Read a comma-separated values (csv) file into DataFrame.

Also supports optionally iterating or breaking of the file into chunks.

Additional help can be found in the online docs for

`IO Tools <[http://pandas.pydata.org/pandas-docs/stable/user\\_guide/i](http://pandas.pydata.org/pandas-docs/stable/user_guide/i)o.html>`\_.

Parameters

-----

filepath\_or\_buffer : str, path object or file-like object

Any valid string path is acceptable. The string could be a URL.

Valid

URL schemes include http, ftp, s3, and file. For file URLs, a h  
ost is  
expected. A local file could be: file://localhost/path/to/tabl  
e.csv.

If you want to pass in a path object, pandas accepts any ``os.P  
athLike``.

By file-like object, we refer to objects with a ``read()`` meth  
od, such as

a file handler (e.g. via builtin ``open`` function) or ``String  
IO``.

sep : str, default ',',

Delimiter to use. If sep is None, the C engine cannot automatic  
ally detect

the separator, but the Python parsing engine can, meaning the l  
atter will

be used and automatically detect the separator by Python's buil  
tin sniffer

tool, ``csv.Sniffer``. In addition, separators longer than 1 ch  
aracter and

different from ``'\s+'`` will be interpreted as regular express  
ions and

will also force the use of the Python parsing engine. Note that  
regex

delimiters are prone to ignoring quoted data. Regex example: `  
``'\r\t'``.

delimiter : str, default ``None``

Alias for sep.

header : int, list of int, default 'infer'

Row number(s) to use as the column names, and the start of the data. Default behavior is to infer the column names: if no names are passed the behavior is identical to ``header=0`` and column names are inferred from the first line of the file, if column names are passed explicitly then the behavior is identical to ``header=None``. Explicitly pass ``header=0`` to be able to replace existing names. The header can be a list of integers that specify row locations for a multi-index on the columns e.g. [0,1,3]. Intervening rows that are not specified will be skipped (e.g. 2 in this example is skipped). Note that this parameter ignores commented lines and empty lines if ``skip\_blank\_lines=True``, so ``header=0`` denotes the first line of data rather than the first line of the file.

names : array-like, optional

List of column names to use. If file contains no header row, then you should explicitly pass ``header=None``. Duplicates in this list are not allowed.

index\_col : int, str, sequence of int / str, or False, default ``None``

Column(s) to use as the row labels of the ``DataFrame``, either given as string name or column index. If a sequence of int / str is given, a MultiIndex is used.

Note: ``index\_col=False`` can be used to force pandas to *not* use the first column as the index, e.g. when you have a malformed file with delimiters at the end of each line.

usecols : list-like or callable, optional

Return a subset of the columns. If list-like, all elements must either be positional (i.e. integer indices into the document columns) or strings that correspond to column names provided either by the user in ``names`` or inferred from the document header row(s). For example, a valid list-like ``usecols`` parameter would be ``[0, 1, 2]`` or ``['foo', 'bar', 'baz']``. Element order is ignored, so ``usecols=[0, 1]`` is the same as ``[1, 0]``.

To instantiate a DataFrame from ``data`` with element order preserved use

```
pd.read_csv(data, usecols=['foo', 'bar'])[['foo', 'bar']]
```

for columns in ``['foo', 'bar']`` order or

```
pd.read_csv(data, usecols=['foo', 'bar'])[['bar', 'foo']]
```

for ``['bar', 'foo']`` order.

If callable, the callable function will be evaluated against the column names, returning names where the callable function evaluates to True. An example of a valid callable argument would be ``lambda x: x.upper() in ['AAA', 'BBB', 'DDD']``. Using this parameter results in much faster parsing time and lower memory usage.

`squeeze` : bool, default False  
If the parsed data only contains one column then return a Series.

`prefix` : str, optional  
Prefix to add to column numbers when no header, e.g. 'X' for X0, X1, ...

`mangle_dupe_cols` : bool, default True  
Duplicate columns will be specified as 'X', 'X.1', ...'X.N', rather than 'X'...'X'. Passing in False will cause data to be overwritten if there are duplicate names in the columns.

`dtype` : Type name or dict of column -> type, optional  
Data type for data or columns. E.g. {'a': np.float64, 'b': np.int32, 'c': 'Int64'}  
Use `str` or `object` together with suitable `na\_values` settings to preserve and not interpret dtype.  
If converters are specified, they will be applied INSTEAD of dtype conversion.

`engine` : {'c', 'python'}, optional  
Parser engine to use. The C engine is faster while the python engine is currently more feature-complete.

`converters` : dict, optional  
Dict of functions for converting values in certain columns. Keys can either be integers or column labels.

`true_values` : list, optional  
Values to consider as True.

`false_values` : list, optional  
Values to consider as False.

`skipinitialspace` : bool, default False  
Skip spaces after delimiter.

`skiprows` : list-like, int or callable, optional  
Line numbers to skip (0-indexed) or number of lines to skip (int) at the start of the file.

If callable, the callable function will be evaluated against the row indices, returning True if the row should be skipped and False otherwise.

An example of a valid callable argument would be ``lambda x: x in [0, 2]``.

`skipfooter` : int, default 0

Number of lines at bottom of file to skip (Unsupported with engine='c').

nrows : int, optional  
Number of rows of file to read. Useful for reading pieces of large files.

na\_values : scalar, str, list-like, or dict, optional  
Additional strings to recognize as NA/NaN. If dict passed, specific per-column NA values. By default the following values are interpreted as

NaN: '', '#N/A', '#N/A N/A', '#NA', '-1.#IND', '-1.#QNAN', '-NaN', '-nan',  
'1.#IND', '1.#QNAN', 'N/A', 'NA', 'NULL', 'NaN', 'n/a', 'nan', 'null'.

keep\_default\_na : bool, default True  
Whether or not to include the default NaN values when parsing the data.

Depending on whether `na\_values` is passed in, the behavior is as follows:

- \* If `keep\_default\_na` is True, and `na\_values` are specified, `na\_values` is appended to the default NaN values used for parsing.
- \* If `keep\_default\_na` is True, and `na\_values` are not specified, only the default NaN values are used for parsing.
- \* If `keep\_default\_na` is False, and `na\_values` are specified, only the NaN values specified `na\_values` are used for parsing.
- \* If `keep\_default\_na` is False, and `na\_values` are not specified, no strings will be parsed as NaN.

Note that if `na\_filter` is passed in as False, the `keep\_default\_na` and `na\_values` parameters will be ignored.

na\_filter : bool, default True  
Detect missing value markers (empty strings and the value of na\_values). In data without any NAs, passing na\_filter=False can improve the performance of reading a large file.

verbose : bool, default False  
Indicate number of NA values placed in non-numeric columns.

skip\_blank\_lines : bool, default True  
If True, skip over blank lines rather than interpreting as NaN values.

parse\_dates : bool or list of int or names or list of lists or dict, default False

The behavior is as follows:

- \* boolean. If True -> try parsing the index.
- \* list of int or names. e.g. If [1, 2, 3] -> try parsing columns 1, 2, 3 each as a separate date column.
- \* list of lists. e.g. If [[1, 3]] -> combine columns 1 and 3 and parse as



```

    a single date column.
    * dict, e.g. {'foo' : [1, 3]} -> parse columns 1, 3 as date and
call
    result 'foo'

    If a column or index cannot be represented as an array of datet
imes,
    say because of an unparseable value or a mixture of timezones,
the column
    or index will be returned unaltered as an object data type. For
non-standard datetime parsing, use ``pd.to_datetime`` after
``pd.read_csv``. To parse an index or column with a mixture of
timezones,
    specify ``date_parser`` to be a partially-applied
:func:`pandas.to_datetime` with ``utc=True``. See
:ref:`io.csv.mixed_timezones` for more.

    Note: A fast-path exists for iso8601-formatted dates.
infer_datetime_format : bool, default False
    If True and `parse_dates` is enabled, pandas will attempt to in
fer the
    format of the datetime strings in the columns, and if it can be
inferred,
    switch to a faster method of parsing them. In some cases this c
an increase
    the parsing speed by 5-10x.
keep_date_col : bool, default False
    If True and `parse_dates` specifies combining multiple columns
then
    keep the original columns.
date_parser : function, optional
    Function to use for converting a sequence of string columns to
an array of
    datetime instances. The default uses ``dateutil.parser.parser``
to do the
    conversion. Pandas will try to call `date_parser` in three diff
erent ways,
    advancing to the next if an exception occurs: 1) Pass one or mo
re arrays
    (as defined by `parse_dates`) as arguments; 2) concatenate (row
-wise) the
    string values from the columns defined by `parse_dates` into a
single array
    and pass that; and 3) call `date_parser` once for each row usin
g one or
    more strings (corresponding to the columns defined by `parse_da
tes`) as
    arguments.
dayfirst : bool, default False
    DD/MM format dates, international and European format.
cache_dates : boolean, default True
    If True, use a cache of unique, converted dates to apply the da
atetime
    conversion. May produce significant speed-up when parsing dupli
cate
    date strings, especially ones with timezone offsets.

```

```

        .. versionadded:: 0.25.0
    iterator : bool, default False
        Return TextFileReader object for iteration or getting chunks with
        ``get_chunk()``.
    chunksize : int, optional
        Return TextFileReader object for iteration.
        See the `IO Tools docs
        <http://pandas.pydata.org/pandas-docs/stable/io.html#io-chunking>`_
        for more information on ``iterator`` and ``chunksize``.
    compression : {'infer', 'gzip', 'bz2', 'zip', 'xz', None}, default 'infer'
        For on-the-fly decompression of on-disk data. If 'infer' and
        `filepath_or_buffer` is path-like, then detect compression from
        the following extensions: '.gz', '.bz2', '.zip', or '.xz' (otherwise no
        decompression). If using 'zip', the ZIP file must contain only one data
        file to be read in. Set to None for no decompression.

        .. versionadded:: 0.18.1 support for 'zip' and 'xz' compression.

    thousands : str, optional
        Thousands separator.
    decimal : str, default '.'
        Character to recognize as decimal point (e.g. use ',' for European data).
    lineterminator : str (length 1), optional
        Character to break file into lines. Only valid with C parser.
    quotechar : str (length 1), optional
        The character used to denote the start and end of a quoted item. Quoted
        items can include the delimiter and it will be ignored.
    quoting : int or csv.QUOTE_* instance, default 0
        Control field quoting behavior per `csv.QUOTE_*` constants. Use one of
        QUOTE_MINIMAL (0), QUOTE_ALL (1), QUOTE_NONNUMERIC (2) or QUOTE_NONE (3).
    doublequote : bool, default ``True``
        When quotechar is specified and quoting is not `QUOTE_NONE`, indicate
        whether or not to interpret two consecutive quotechar elements INSIDE a
        field as a single `quotechar` element.
    escapechar : str (length 1), optional
        One-character string used to escape other characters.
    comment : str, optional
        Indicates remainder of line should not be parsed. If found at the beginning
        of a line, the line will be ignored altogether. This parameter must be a
        single character. Like empty lines (as long as `skip_blank_lines=True`),
        fully commented lines are ignored by the parameter `header` but

```

not by  
``skiprows``. For example, if ``comment='#'``, parsing  
``#empty\na,b,c\n1,2,3`` with ``header=0`` will result in 'a,b,  
c' being  
treated as the header.  
encoding : str, optional  
Encoding to use for UTF when reading/writing (ex. 'utf-8'). `Li  
st of Python  
standard encodings  
<<https://docs.python.org/3/library/codecs.html#standard-encodin>  
gs>`\_` .  
dialect : str or csv.Dialect, optional  
If provided, this parameter will override values (default or no  
t) for the  
following parameters: ``delimiter``, ``doublequote``, ``escapechar``,  
``skipinitialspace``, ``quotechar``, and ``quoting``. If it is necess  
ary to  
override values, a `ParserWarning` will be issued. See csv.Dialec  
t  
documentation for more details.  
error\_bad\_lines : bool, default True  
Lines with too many fields (e.g. a csv line with too many comma  
s) will by  
default cause an exception to be raised, and no `DataFrame` will  
be returned.  
If False, then these "bad lines" will dropped from the `DataFrame`  
e that is  
returned.  
warn\_bad\_lines : bool, default True  
If `error_bad_lines` is False, and `warn_bad_lines` is True, a warn  
ing for each  
"bad line" will be output.  
delim\_whitespace : bool, default False  
Specifies whether or not whitespace (e.g. `` `` or `` ``  
`) will be  
used as the sep. Equivalent to setting ``sep='\s+'``. If this o  
ption  
is set to True, nothing should be passed in for the ``delimiter``  
parameter.  
.. versionadded:: 0.18.1 support for the Python parser.  
low\_memory : bool, default True  
Internally process the file in chunks, resulting in lower memor  
y use  
while parsing, but possibly mixed type inference. To ensure no  
mixed  
types either set False, or specify the type with the ``dtype`` pa  
rameter.  
Note that the entire file is read into a single `DataFrame` regar  
dless,  
use the ``chunksize`` or ``iterator`` parameter to return the data  
in chunks.  
(Only valid with C parser).  
memory\_map : bool, default False  
If a filepath is provided for ``filepath_or_buffer``, map the fil

e object  
 directly onto memory and access the data directly from there. Using this  
 option can improve performance because there is no longer any I/O overhead.

`float_precision` : str, optional  
 Specifies which converter the C engine should use for floating-point  
 values. The options are ``None`` for the ordinary converter, ``high`` for the high-precision converter, and ``round_trip`` for the  
 round-trip converter.

Returns  
 -----  
 DataFrame or TextParser  
 A comma-separated values (csv) file is returned as two-dimensional  
 data structure with labeled axes.

See Also  
 -----  
`to_csv` : Write DataFrame to a comma-separated values (csv) file.  
`read_csv` : Read a comma-separated values (csv) file into DataFrame.  
`read_fwf` : Read a table of fixed-width formatted lines into DataFrame.

Examples  
 -----  
 >>> pd.read\_csv('data.csv') # doctest: +SKIP

## Exercise 1

How should we read in `adult_test.csv` properly? Identify and fix the problem.

```
In [3]: df = pd.read_csv('data/adult_test.csv')
```

```
print(df.head())
```

```
      This is the test set for the adult dataset.  Unnamed: 1  Unnamed: 2  \
0      The first two lines need to be skipped.      NaN      NaN
1                                     age  workclass  fnlwgt
2                                     25   Private  226802
3                                     38   Private  89814
4                                     28  Local-gov  336951

      Unnamed: 3      Unnamed: 4      Unnamed: 5      Unnamed: 6
\
0      NaN      NaN      NaN      NaN
1  education  education-num  marital-status  occupation
2      11th      7      Never-married  Machine-op-inspct
3   HS-grad      9  Married-civ-spouse  Farming-fishing
4  Assoc-acdm      12  Married-civ-spouse  Protective-serv

      Unnamed: 7  Unnamed: 8  Unnamed: 9  Unnamed: 10  Unnamed: 11  \
0      NaN      NaN      NaN      NaN      NaN
1  relationship      race      sex  capital-gain  capital-loss
2   Own-child      Black      Male      0      0
3   Husband      White      Male      0      0
4   Husband      White      Male      0      0

      Unnamed: 12      Unnamed: 13  Unnamed: 14
0      NaN      NaN      NaN
1  hours-per-week  native-country  gross-income
2      40  United-States  <=50K.
3      50  United-States  <=50K.
4      40  United-States  >50K.
```

```
In [4]: # two solutions
df = pd.read_csv('data/adult_test.csv',header=2)
df = pd.read_csv('data/adult_test.csv',skiprows=2)
print(df.head())
```

	age	workclass	fnlwgt	education	education-num	marital-
0	25	Private	226802	11th	7	Never-m
1	38	Private	89814	HS-grad	9	Married-civ-
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-
3	44	Private	160323	Some-college	10	Married-civ-
4	18	?	103497	Some-college	10	Never-m

  

	occupation	relationship	race	sex	capital-gain \
0	Machine-op-inspct	Own-child	Black	Male	0
1	Farming-fishing	Husband	White	Male	0
2	Protective-serv	Husband	White	Male	0
3	Machine-op-inspct	Husband	Black	Male	7688
4	?	Own-child	White	Female	0

  

	capital-loss	hours-per-week	native-country	gross-income
0	0	40	United-States	<=50K.
1	0	50	United-States	<=50K.
2	0	40	United-States	>50K.
3	0	40	United-States	>50K.
4	0	30	United-States	<=50K.

## Data transformations: pandas data frames

By the end of this talk, you will be able to

- read in csv, excel, and sql data into a pandas data frame
- **filter rows in various ways**
- select columns
- merge and append data frames

## How to select rows?

*1) Integer-based indexing, numpy arrays are indexed the same way.*

*2) Select rows based on the value of the index column*

*3) select rows based on column condition*

### 1) Integer-based indexing, numpy arrays are indexed the same way.

```
In [5]: # df.iloc[] - for more info, see https://pandas.pydata.org/pandas-docs/s
table/user_guide/indexing.html#indexing-integer
# iloc is how numpy arrays are indexed (non-standard python indexing)

# [start:stop:step] - general indexing format

# start stop step are optional
#print(df.iloc[:])
#print(df.iloc[:,])
#print(df.iloc[:,1])

# select one row - 0-based indexing
#print(df.iloc[3])

# indexing from the end of the data frame
#print(df.iloc[-2])
```

```
In [6]: # select a slice - stop index not included
#print(df.iloc[3:7])

# select every second element of the slice - stop index not included
# print(df.iloc[3:7:2])

#print(df.iloc[3:7:-2]) # return empty dataframe
#print(df.iloc[7:3:-2])# return rows with indices 7 and 5. 3 is the sto
p so it is not included

# can be used to reverse rows
#print(df.iloc[::-1])

# here is where indexing gets non-standard python
# select the 2nd, 5th, and 10th rows
#print(df.iloc[[1,4,9]]) # such indexing doesn't work with lists but it
works with numpy arrays
```

### 2) Select rows based on the value of the index column

```
In [7]: # df.loc[] - for more info, see https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#indexing-label

# print(df.index) # the default index when reading in a file is a range index. In this case,
#               # .loc and .iloc works ALMOST the same.
# one difference:
# print(df.loc[3:9:2]) # this selects the 4th, 6th, 8th, 10th rows - the stop element is included!

# help(df.set_index)
```

```
In [8]: # df_index_age = df.set_index('age', drop=False)

# print(df_index_age.index)
# print(df_index_age.head())

# print(df_index_age.loc[30].head()) # collect everyone with age 30 - the index is non-unique

# print(df_index_age.loc[30:35]) # non-default index cannot be sliced.
#                               # this does not return everyone between ages of 30 and 35
```

### 3) select rows based on column condition

```
In [9]: # one condition
# print(df[df['age']==30].head())
# here is the condition: it's a boolean series - series is basically a dataframe with one column
# print(df['age']==30)

# multiple conditions can be combined with & (and) | (or)
# print(df[(df['age']>30)&(df['age']<35)].head())
# print(df[(df['age']==90)|(df['native-country']==' Hungary')])
```

## Exercise 2

How many people in adult\_data.csv work at least 60 hours a week and have a doctorate?



```
In [10]: # solution
df = pd.read_csv('data/adult_data.csv')
print(df[(df['hours-per-week'] >= 60)&(df['education']==' Doctorate')].shape[0])
# [96 rows x 15 columns]
# we will learn how to modify columns and remove irregularities like this later.

# mistakes the students could make:
print(df[(df['hours-per-week'] >= 60)&(df['education']==' Doctorate')].shape)

print(df[(df['hours-per-week'] > 60)&(df['education']==' Doctorate')].shape)

df = pd.read_csv('data/adult_test.csv', skiprows=2)
print(df[(df['hours-per-week'] >= 60)&(df['education']==' Doctorate')].shape)

print(df[(df['hours-per-week'] > 60)&(df['education']==' Doctorate')].shape)

# these are all good possibilities for tophat answers
```

```
96
(0, 15)
(39, 15)
(33, 15)
(11, 15)
```

## Data transformations: pandas data frames

By the end of this talk, you will be able to

- read in csv, excel, and sql data into a pandas data frame
- filter rows in various ways
- **select columns**
- merge and append data frames

```
In [11]: columns = df.columns
print(columns)

# select columns by column name
#print(df[['age', 'hours-per-week']])
#print(columns[[1,5,7]])
#print(df[columns[[1,5,7]]])

# select columns by index using iloc
#print(df.iloc[:,3])

# select columns by index - not standard python indexing
#print(df.iloc[:,[3,5,6]])

# select columns by index - standard python indexing
#print(df.iloc[:, :2])

Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',
       'marital-status', 'occupation', 'relationship', 'race', 'sex',
       'capital-gain', 'capital-loss', 'hours-per-week', 'native-country',
       'gross-income'],
      dtype='object')
```

## Data transformations: pandas data frames

By the end of this talk, you will be able to

- read in csv, excel, and sql data into a pandas data frame
- filter rows in various ways
- select columns
- **merge and append data frames**

### How to merge dataframes?

Merge - info on data points are distributed in multiple files

```
In [12]: # We have two datasets from two hospitals

hospital1 = {'ID':['ID1','ID2','ID3','ID4','ID5','ID6','ID7'],'col1':[5,
8,2,6,0,2,5],'col2':['y','j','w','b','a','b','t']}
df1 = pd.DataFrame(data=hospital1)
print(df1)

hospital2 = {'ID':['ID2','ID5','ID6','ID10','ID11'],'col3':[12,76,34,98,
65],'col2':['q','u','e','l','p']}
df2 = pd.DataFrame(data=hospital2)
print(df2)
```

	ID	col1	col2
0	ID1	5	y
1	ID2	8	j
2	ID3	2	w
3	ID4	6	b
4	ID5	0	a
5	ID6	2	b
6	ID7	5	t

  

	ID	col3	col2
0	ID2	12	q
1	ID5	76	u
2	ID6	34	e
3	ID10	98	l
4	ID11	65	p

```
In [13]: # we are interested in only patients from hospital1
# df_left = df1.merge(df2,how='left',on='ID') # IDs from the left dataframe (df1) are kept
# print(df_left)

# we are interested in only patients from hospital2
# df_right = df1.merge(df2,how='right',on='ID') # IDs from the right dataframe (df2) are kept
# print(df_right)

# we are interested in patients who were in both hospitals
# df_inner = df1.merge(df2,how='inner',on='ID') # merging on IDs present in both dataframes
# print(df_inner)

# we are interested in all patients who visited at least one of the hospitals
# df_outer = df1.merge(df2,how='outer',on='ID') # merging on IDs present in any dataframe
# print(df_outer)
```

## How to append dataframes?

Append - new data comes in over a period of time. E.g., one file per month/quarter/fiscal year etc.

You want to combine these files into one data frame.

```
In [14]: df_append = df1.append(df2) # note that rows with ID2, ID5, and ID6 are
        duplicated! Indices are duplicated too.
        print(df_append)

        # df_append = df1.append(df2,ignore_index=True) # note that rows with ID
        2, ID5, and ID6 are duplicated!
        # print(df_append)

        # d3 = {'ID':['ID23','ID94','ID56','ID17'],'col1':['rt','h','st','n
        e'],'col2':[23,86,23,78]}
        # df3 = pd.DataFrame(data=d3)
        # print(df3)

        # df_append = df1.append([df2,df3],ignore_index=True) # multiple datafra
        mes can be appended to df1
        # print(df_append)
```

	ID	col1	col2	col3
0	ID1	5.0	y	NaN
1	ID2	8.0	j	NaN
2	ID3	2.0	w	NaN
3	ID4	6.0	b	NaN
4	ID5	0.0	a	NaN
5	ID6	2.0	b	NaN
6	ID7	5.0	t	NaN
0	ID2	NaN	q	12.0
1	ID5	NaN	u	76.0
2	ID6	NaN	e	34.0
3	ID10	NaN	l	98.0
4	ID11	NaN	p	65.0

/anaconda3/envs/datasci\_v0.0.2\_local4/lib/python3.6/site-packages/pandas/core/frame.py:7138: FutureWarning: Sorting because non-concatenation axis is not aligned. A future version of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=True'.

sort=sort,

## Exercise 3

```
In [15]: raw_data_1 = {
    'subject_id': ['1', '2', '3', '4', '5'],
    'first_name': ['Alex', 'Amy', 'Allen', 'Alice', 'Ayoung'],
    'last_name': ['Anderson', 'Ackerman', 'Ali', 'Aoni', 'Atiches']}

raw_data_2 = {
    'subject_id': ['6', '7', '8', '9', '10'],
    'first_name': ['Billy', 'Brian', 'Bran', 'Bryce', 'Betty'],
    'last_name': ['Bonder', 'Black', 'Balwner', 'Brice', 'Btisan']}

raw_data_3 = {
    'subject_id': ['1', '2', '3', '4', '5', '7', '8', '9', '10', '11'],
    'test_id': [51, 15, 15, 61, 16, 14, 15, 1, 61, 16]}

# Create three data frames from raw_data_1, 2, and 3.
# Append the first two data frames and assign it to df_append.
# Merge the third data frame with df_append such that only subject_ids from df_append are present.
# Assign the new data frame to df_merge.
# How many rows and columns do we have in df_merge?
```

```
In [16]: # The solution
```

```
df1 = pd.DataFrame(raw_data_1)
df2 = pd.DataFrame(raw_data_2)
df3 = pd.DataFrame(raw_data_3)

df_append = df1.append(df2)
print(df_append)

df_merge = df_append.merge(df3, how='left', on='subject_id')
print(df_merge)
print(df_merge.shape)
```

	subject_id	first_name	last_name
0	1	Alex	Anderson
1	2	Amy	Ackerman
2	3	Allen	Ali
3	4	Alice	Aoni
4	5	Ayoung	Atiches
0	6	Billy	Bonder
1	7	Brian	Black
2	8	Bran	Balwner
3	9	Bryce	Brice
4	10	Betty	Btisan

  

	subject_id	first_name	last_name	test_id
0	1	Alex	Anderson	51.0
1	2	Amy	Ackerman	15.0
2	3	Allen	Ali	15.0
3	4	Alice	Aoni	61.0
4	5	Ayoung	Atiches	16.0
5	6	Billy	Bonder	NaN
6	7	Brian	Black	14.0
7	8	Bran	Balwner	15.0
8	9	Bryce	Brice	1.0
9	10	Betty	Btisan	61.0

(10, 4)

**Always check that the resulting dataframe is what you wanted to end up with!**

- small toy datasets are ideal to test your code.

**If you need to do a more complicated dataframe operation, check out `pd.concat()`!**

**We will learn how to add/delete/modify columns later when we learn about feature engineering.**

## By now, you are able to

- read in csv, excel, and sql data into a pandas data frame
- filter rows in various ways
- select columns
- merge and append data frames

In [ ]: